Fisher Divergence for mass matrix adaptation

Adrian Seyboldt

October 5, 2022

1 Introduction

The choice of mass matrix is known to be critical for the performance of HMC samplers. Mass matrices are usually adapted during tuning, and convergence speed to a good mass matrix has a big impact on the total cost of posterior sampling as well.

todo add references to important hmc/nuts papers.

In the case of normal posterior $N(\mu, \Sigma)$ we can use the condition number of Σ as an estimate of how difficult it is to sample from this posterior with an identity mass matrix.

$$\kappa(\Sigma) = \frac{\max \text{eigvals}(\Sigma)}{\min \text{eigvals}(\Sigma)}$$

If we use a mass matrix $\hat{\Sigma}^{-1}$, we can compute the condition number by replacing the eigenvalues by the generalized eigenvalues with respect to the inverse mass matrix eigvals$(\Sigma, \hat{\Sigma})$.

A better estimate can be found here[1], but for our purposes the simpler condition number will do fine.

Currently, most implementations (at least Stan + PyMC) use a regularized version of the empirical covariance or its diagonal as inverse mass matrix.

1.1 Use gradients of posterior log density

This means however, that we do not use an additional source of information that we have available anyway: the gradients of the posterior log density at the positions of the draws.

Since the gradients transform covariantly, where the draws transform contravariantly, the covariance of gradients is the inverse covariance matrix for a Gaussian posterior. This indicates that the covariance of the gradients gives us an estimate for Σ^{-1}, while the variance of the draws themselves gives an estimate for Σ. If we have few draws compared to the dimensionality of the problem, the empirical covariances therefore contain information about the large and small eigenvalues of the true covariance of the posterior respectively. But since both of those are important for a good mass matrix, as can be seen from the condition number, it stands to reason that we can use the gradient information to improve mass matrix estimates.

---

[1] https://arxiv.org/pdf/1905.09813.pdf

Estimates of the variance with only the draws are also limited by the Cramér-Rao bound, which is no longer the case if we consider additional information. For instance, for a 1D Gaussian posterior one draw with its gradient is already enough to uniquely identify the posterior distribution, and as we will see it also provides a lot of information about the posterior in the higher dimensional case.

1.2 Fisher Divergence

Choosing a mass matrix is equivalent to fitting a Gaussian distribution to the posterior, where the inverse covariance matrix of the fitted distribution is the mass matrix.

A natural way to incorporate the gradient information into this covariance estimation is by minimizing the Fisher divergence (see e.g. https://arxiv.org/pdf/1905.05284.pdf) between our estimated normal distribution with density q and the actual posterior distribution with density p:

$$F(p, q) = \int \|\nabla \log p(\theta) - \nabla \log q(\theta)\|^2 p(\theta) d\theta.$$

It is not immediately obvious which norm we should use here. A natural choice becomes apparent if we remember that the mass matrix as used in HMC itself defines a norm on the tangent parameter space. I propose to use this norm, which means that we are looking for a distribution with density q who's gradients are close to the gradients of the target distribution p *as seen by the sampler*:

$$F(p, q) = \int \|\nabla \log p(\theta) - \nabla \log q(\theta)\|_{\hat{\Sigma}}^2 p(\theta) d\theta,$$

where the norm on the tangent space is given by $\|x\|_{\hat{\Sigma}}^2 = x^T \hat{\Sigma} x$, where $\hat{\Sigma}$ is the covariance of our fitted distribution q.

This is equivalent to finding an affine transformation $\phi(x) = Ax + \mu$ such that

$$F(\phi) = \int \|\nabla \log p_\phi(\theta) - \nabla \log N(\phi(\theta) \mid 0, 1)\|^2 p(\theta) d\theta \tag{1}$$

$$= \int \|\nabla \log p_\phi(\theta) + \phi(\theta)\|^2 p(\theta) d\theta \tag{2}$$

is minimal, where $p_\phi(\theta) = |J_\phi|^{1/2} p(\phi(\theta))$ is the transformed density and $AA^T = \hat{\Sigma}$ corresponds to the covariance of our fitted Gaussian distribution, or the inverse mass matrix and ϕ^* is the pullback of ϕ.

Proof. TODO □

This definition is now valid for arbitrary families of diffeomorphisms ϕ.

During sampling we have N draws θ_i and gradients $\nabla \log p(\theta_i)$, so we approximate F by its Monte Carlo estimate:

$$\hat{F}(p, q) = \frac{1}{N} \sum_i \|\nabla \log p(\theta_i) - \nabla \log q(\theta_i)\|_\Sigma^2$$

or

$$\hat{F}(\phi) = \frac{1}{N} \sum_i \|\nabla \log p_\phi(\theta_i) + \phi(\theta_i)\|^2$$

1.3 Properties of minimization problem

Theorem 1. *If $\phi_{\mu,A} = Ax + \mu$ and X with density p is such that $\mathrm{Cov}(X)$ and $\mathrm{Cov}(\nabla \log(p(X)))$ exist and are positive definite, then $F(\phi_{\mu,A})$ is minimal iff $\mu = \mathrm{E}(X)$ and AA^T is the geodesic mean of $\mathrm{Cov}(X)$ and $\mathrm{Cov}(\nabla \log(p(X)))^{-1}$.*

Put differently, AA^T is such that

$$d(\mathrm{Cov}(X), AA^T) + d(AA^T, \mathrm{Cov}(\nabla \log(p(X)))^{-1}),$$

is minimal, where $d(\Sigma, \Omega) = \|\log(\mathrm{eigvals}(\Sigma, \Omega))\|$ is the geodesic distance of symmetric positive definite matrices Σ and Ω with respect to the intrinsic metric.

$\hat{F}$ is minimal for all A with $\hat{\Sigma} A A^T \hat{\Sigma} = \hat{\Omega}$, where $\hat{\Sigma} = \mathrm{Cov}(X_i)$ is the empirical covariance of the draws and $\hat{\Omega} = \mathrm{Cov}(\nabla \log(p(X_i)))$ is the empirical covariance of the gradients.

Proof. See appendix. $\qquad\square$

So if we interpret both $\mathrm{Cov}(X)$ and $\mathrm{Cov}(\nabla \log(p(X)))^{-1}$ as possible choices for the inverse mass matrix, minimizing F leads to the mean of those two choices. If we have more dimensions than draws, this minimum (or mean) is not unique however, so we will need regularization of some kind in order to find a unique mass matrix estimate.

Corollary 1. *If ϕ is as above and additionally $X \sim N(\mu', \Sigma')$, then the minimum of F is at $\mu = \mu'$ and $\Sigma = \Sigma'$, because in this case $\mathrm{Cov}(X) = \mathrm{Cov}(\nabla \log(p))^{-1}$.*

Let k be the number of draws in n dimensions and let $\hat{A}$ be the linear part of the minimizer of $\hat{F}$. Then there exists a k-dimensional subspace of $\mathrm{span}(X_i, \nabla \log p(X_i))$ in which $\hat{A}\hat{A}^T$ matches the true covariance Σ exactly.

This means if our posterior is a Gaussian, we get better estimates for the mass matrix the more draws we have, and if we have more draws than dimensions our mass matrix estimate will be perfect.

Corollary 2. *Given a space of transformation $\phi_{\mu,\sigma}(X) = \mu + \sigma \odot X$ (this corresponds to diagonal mass matrices) F is minimal if $\mu = E_p[X]$ and $\sigma^2 = \sqrt{\frac{\mathrm{Var}(X)}{\mathrm{Var}(\nabla \log(p(X)))}}$ and*

$$\sigma = \arg\min_{\sigma} d(\mathrm{Cov}(X), \mathrm{diag}(\sigma)^2) + d(\mathrm{diag}(\sigma)^2, \mathrm{Cov}(\nabla \log(p(X)))^{-1}),$$

$\hat{F}$ is minimal if

$$\sigma^2 = \sqrt{\frac{\mathrm{Var}(X_i)}{\mathrm{Var}(\nabla \log(p(X_i)))}}.$$

This last result motivates a simple modification of the status quo diagonal mass matrix adaptation. Instead of adapting based on the marginal variance of each parameter, we adapt based on the geometric mean of the marginal variance and inverse marginal variance of the gradients.

In the next part of this paper we will investigate this modification in more detail, and compare an implementation of it in nutpie with the current default choice.

In later parts we will show how we can develop the ideas further for sub-quadratic mass matrix adaptation that still can deal with highly correlated posteriors.

The final part will investigate further generalizations to non-linear transformations using neural networks.

2 Diagonal mass matrix

Based on the previous results we can estimate diagonal mass matrices as

$$\hat{\Sigma}^{-1} = \text{diag}\left(\sqrt{\frac{\text{Var}(X_i)}{\text{Var}(\nabla \log(p(X_i)))}}\right).$$

We have seen that if the posterior is Gaussian with diagonal covariance, this will recover the exact covariance matrix with only two draws.

Posterior distributions often approximate diagonal covariances, but have some non-diagonal components. To evaluate the new proposed diagonal mass matrix approximation for non-diagonal Gaussian distributions we model posterior covariance matrices as coming from a distribution defined as follows (todo add ref for Haar distribution, uniform on orthogonal matrices):

$$U \sim \text{Haar}(n)$$
$$\log(\lambda) \sim N(0, \sigma_{\text{vals}}^2)$$
$$\log(d) \sim N(0, \sigma_{diag}^2)$$
$$\Sigma = \text{diag}(d) U \text{diag}(\lambda) U^T \text{diag}(d)$$

We have two hyperparameters, σ_{vals} and σ_{diag}. The first controls how much correlation structure we have (if $\sigma_{\text{vals}} = 0$ the correlation matrix is diagonal), the second how much the marginal variables are scaled differently. If we sample from this distribution of covariance matrices we can apply both mass matrix adaptation methods and compare their respective quality using the condition number from (todo add ref). Figure 1 shows the results of this comparison, and we can see that the Fisher divergence based adaptation is clearly superior if we have few draws or if the matrix is dominated by the diagonal scaling. I have not found any covariance matrices so far where the fisher divergence based estimate was worse (I guess that doesn't exist, but I don't have a proof).

todo:

- Show faster convergence during early tuning

- Show ess/grad in a few real models

- Overall better performance of nutpie compared to stan/PyMC sampler implementations

3 Mass matrix with correlation structure

Discuss some simple regularizers (is $tr(\Sigma) + tr(\Omega)$), even though they don't seem to work all that well?

Discuss projection onto $\text{span}(X, \nabla X)$, to reduce from $O(n^2)$ to $O(k^2)$.

Large influence of diagonal in many actual models, i.e. scale draws first?

Add epsilon to 0 eigvals trick?

Finally, the one implemented in covadapt: Matrix representation as $D(U(\Lambda - I)U^T + I)D$, so that we can avoid strong regularization on D but have regularization on Λ, and Λ can be much smaller than $n \times n$, so we avoid quadratic cost: only fix a few eigenvector, eigenvalue combinations, that cause the biggest trouble.
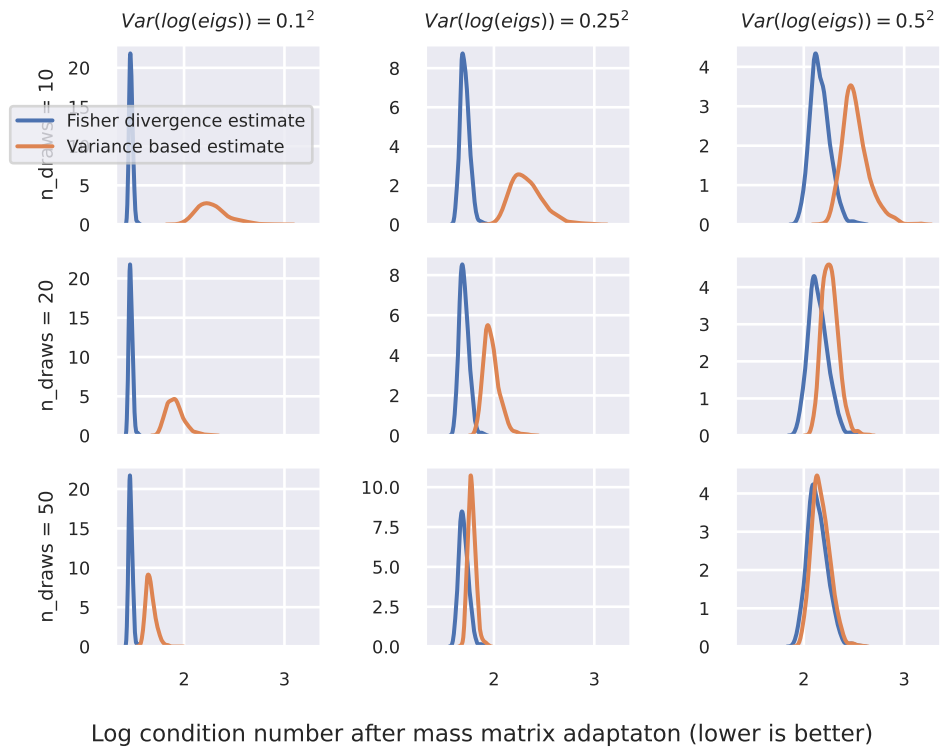
Figure 1: Comparison of fisher divergence based diagonal mass matrix adaptation with variance based

But we actually have to solve the minimizing problem, where we use natural gradient decent. The parameter space is $\log(d), U, \log(\lambda) \in \mathbb{R}^n \times \text{Stiefel}^k \times \mathbb{R}^k$. We have a natural metric, because each point in parameter space refers to a probability distribution (information geometry, fisher metric). We can find gradient with respect to that metric and run manifold minimization.

The resulting algorithm: Initialize mass matrix, draw a few samples, run a few optimization iterations and update mass matrix, repeat...

Show results from a couple of models (e.g. GP where we often have correlations).

4 Non-linear transformations

We need to parameterize transformations ϕ. There is some literature on how to do that with neural networks, I don't know that too well though.

Came up with this so far (which does seem to work quite well?):

Use a bijective neural network directly as transformation.

Activation function based on $\text{act}(x) = (1 - \alpha)\log(1 + x) + \alpha x$ for positive, inverse (based on lambert W function) for negative values of x. This leads to nice nonlinearity, but avoids overflows and underflows.

Each layer is $\text{act}(DA_v x + b)$, where D diagonal, and $A_v = I + (1 - \|v\|^{-2})vv^T$, such that all eigenvalues are one except one with eigenvector v and eigenvalue $\|v\|^2$.

This means that each jacobian determinant is easy to compute, and we can easily compute the objective $\hat{F}$.

Optimization can use SGD or maybe Levenberg–Marquardt with iterative matrix solve using CG or so? Current implementation is a very messy notebook with some jax code...

Appendix

Proof. TODO $\qquad \square$